

Parametric Polymorphism in C++, Java, and C#: Generics vs. Templates and Metaprogramming Power

Prompt Log (Reorganized by Task)

Elia Leonardi
University of Pisa

July 2026

Contents

1	AI Prompt Engineering	2
1.1	Architect Chat	2
1.2	Prompt Engineering Phases	2
1.2.1	Prompt Improvement	2
1.2.2	Prompt Execution	2
1.3	Prompt 0: Configure Prompt Engineering (Task 1 – Language Implementation)	3
1.4	Prompt 0bis: Configure Prompt Engineering (Task 2 – Metaprogramming)	6
2	Task 1 – Language Implementation of Parametric Polymorphism	10
2.1	Prompt 1: Overview of Parametric Polymorphism	11
2.2	Prompt 2: Instantiation Time – C++, Java, and C# Compared . . .	12
2.3	Prompt 3: Runtime Representation – C++, Java, and C# Compared	14
2.4	Prompt 4: Performance Characteristics – C++, Java, and C# Compared	17
2.5	Prompt 5: Type Safety – C++, Java, and C# Compared	19
2.6	Prompt 6: Reflection and Runtime Type Information – C++, Java, and C# Compared	21
3	Task 2 – Metaprogramming Across C++, Java, and C#	23
3.1	Prompt 1: Concrete C++ Template Examples of Compile-Time Computation	23
3.2	Prompt 2: SFINAE and Type Traits	25
3.3	Prompt 3: Variadic Templates, Template Template Parameters, and CRTP	27
3.4	Prompt 4: Limits of Java Generics in Replicating C++ SFINAE and Type Traits	29
3.5	Prompt 5: Limits of C# Generics in Replicating C++ SFINAE and Type Traits	31
3.6	Prompt 6: Instantiation-Time vs. Declaration-Time Checking	33
3.7	Prompt 7: External Metaprogramming Workarounds in Java and C#	36

Chapter 1

AI Prompt Engineering

To ensure the technical accuracy of the generated content and strictly adhere to the project guidelines, this report was developed using a **two-phase prompting strategy**: prompt improvement and prompt execution. A fundamental principle of this methodology is strict context isolation: both the refinement of the prompts and the generation of the final outputs are treated as completely independent operations.

1.1 Architect Chat

An initial “Architect Chat” was utilized to plan and outline the optimal sequence of prompts, ensuring that all project requirements were systematically addressed before any content generation began.

1.2 Prompt Engineering Phases

1.2.1 Prompt Improvement

The first phase focuses on prompt refinement. During this stage, the initial prompts are iteratively improved to align with the project guidelines and maintain a rigorous academic standard. Crucially, the generation of each improved prompt is an independent operation. The AI is tasked solely with rewriting and optimizing the prompt’s structure, without executing it or carrying over context from unrelated topics.

1.2.2 Prompt Execution

The second phase involves the execution of the refined prompts to generate the final output. To ensure the integrity of the results, each execution is treated as a strictly independent operation, performed in an isolated context (e.g., a completely new session). This strict separation guarantees that the generated content relies exclusively on the explicit instructions provided in the finalized prompt, preventing any bias or “context leakage” from the improvement phase. Furthermore, during this execution phase, the AI model is explicitly instructed to provide the source references and documentation used to formulate its responses, ensuring traceability and reliability.

1.3 Prompt 0: Configure Prompt Engineering (Task 1 – Language Implementation)

To ensure the generation of rigorous academic sections, this methodology treats prompt generation and the final report generation as strictly independent operations. The process does not rely on manual, ad-hoc prompts for each language. Instead, it employs a *meta-prompting* approach. A **Master Prompt** was designed whose sole purpose is to act as a "Senior Prompt Engineer" in an isolated session. This system takes a raw specification as input and outputs a formal, optimized execution prompt, which is then used in a completely separate session to generate the final text.

Master Prompt Specification

Below is the complete text of the Master Prompt used to orchestrate the prompt engineering phase for Task 1 (language implementation of parametric polymorphism):

System Persona:

```
Act as a Senior Prompt Engineer specialized in designing rigorous
↳ prompts
for the generation of technical-academic reports in computer science,
programming language theory, compiler design, and software
↳ engineering.
```

Goal:

```
Given, as input, a task specification containing a persona, goal,
requirements, writing-style constraints, output structure, and
↳ possible
reference requirements for producing a technical-academic report,
↳ generate a
single, self-contained master prompt ready to be executed in a new
↳ language
model conversation.
```

```
The generated master prompt must preserve the original intent and
↳ must be
optimized to produce a formally structured academic report with
↳ appropriate
technical depth, logical organization, and scientific writing
↳ standards.
```

Requirements:

Task 1, Extraction and Preservation:

```
Parse the input specification and identify:
- The requested system persona or expert role.
```

- The report objective and title, if provided.
- All explicit and implicit technical requirements.
- Required theoretical concepts, comparisons, analyses, and
 - ↪ evaluations.
- Writing-style constraints.
- Required output format and document structure.
- Requirements related to references, citations, bibliography, or
 - ↪ external academic sources.

Preserve the complete meaning of the original specification.

Do not remove requirements.

Do not introduce new technical requirements.

Do not change the intended scope of the report.

Do not simplify technical constraints.

Task 2, Correction and Structuring:

Correct grammatical errors, ambiguities, inconsistent terminology, formatting problems, and unclear instructions present in the input.

Reorganize the extracted information into a clear and logically

↪ structured

master prompt using the following sections:

- System Persona
- Goal
- Requirements
- Writing Style Constraint
- Output Structure

When reference requirements are present, preserve them explicitly in

↪ the

Requirements section.

Task 3, Academic Report Generation Constraints:

Ensure that the generated master prompt instructs the target language

↪ model

to:

- Produce a complete technical-academic report rather than a short
 - ↪ summary.
- Maintain a formal, objective, and scientific writing style.
- Use precise terminology appropriate for computer science
 - ↪ literature.
- Provide a logical progression from theoretical foundations to
 - ↪ technical

analysis.

- Include comparisons, trade-offs, and implementation details when
 - ↪ required.
- Distinguish clearly between theoretical concepts and implementation mechanisms.
- Avoid unsupported claims and clearly identify assumptions.
- ENFORCE STRICT SCOPE CONTROL: Explicitly instruct the target
 - ↪ language model
 - to confine the generated text strictly to the requested core topic.
 - ↪ The model
 - must actively avoid tangential subjects (e.g., memory footprint,
 - ↪ hardware
 - performance, code bloat, garbage collection overhead, or broader
 - ↪ language
 - comparisons) unless they are explicitly demanded by the input
 - ↪ specification.

Task 4, Reference and Citation Requirements:

If the original specification requires the use of references:

- Require the generated report to rely on authoritative academic or
 - ↪ technicalsources.
- Prefer peer-reviewed publications, academic books, official
 - ↪ languagespecifications, and authoritative technical documentation.
- Require consistent citation formatting.
- Require a final bibliography section containing all cited
 - ↪ resources.
- Avoid fabricated references, unverifiable sources, or unsupported
 - ↪ citations.

Task 5, LaTeX Output Requirements:

If the original specification requires LaTeX output:

- Require complete compilation-ready LaTeX code.
- Preserve a professional academic document structure.
- Include appropriate sections, chapters, mathematical notation,
 - ↪ tables,figures, citations, and bibliography commands when required.
- Ensure compatibility with standard LaTeX compilation workflows.

Writing Style Constraint (STRICT):

Do not execute the input specification.

Do not generate the final report.

Do not answer the original technical question.

The output must contain only the generated master prompt.

Do not include:

- explanations,
- comments,
- editorial notes,
- placeholders,
- analysis of the transformation,
- meta-commentary addressed to the user.

Do not use first-person forms ("I", "we") in the generated master
↪ prompt unless
they are explicitly required by the original specification.

Output Structure:

Return only the generated master prompt.

The generated prompt must be formatted as ready-to-use text and
↪ organized
using the following headings:

System Persona

Goal

Requirements

Writing Style Constraint

Output Structure

No additional text must appear before or after the generated master
↪ prompt.

1.4 Prompt 0bis: Configure Prompt Engineering (Task 2 – Metaprogramming)

For Task 2 (metaprogramming across C++, Java, and C#), a second Master Prompt was used. Its purpose is identical to the one described above, but it is specialized to produce a single *section* of an existing chapter (rather than a standalone report), given one focused instruction as input.

System Persona:

Act as a Senior Prompt Engineer specialized in designing rigorous
→ prompts
for the generation of individual sections within a larger technical-
academic report chapter on programming language theory and software
engineering.

Goal:

Given, as input, a single focused instruction (below, marked as
[INPUT PROMPT]), generate a single, self-contained master prompt
→ ready to
be executed in a new language model conversation. The master prompt
→ must
instruct the target language model to produce ONE SECTION of a larger
chapter -- not a standalone report -- on parametric polymorphism and
metaprogramming across C++, Java, and C#.

[INPUT PROMPT]:
"{{PASTE HERE ONE OF THE 9 PROMPTS}}"

Requirements:

Task 1 -- Extraction and Preservation:

Parse [INPUT PROMPT] and identify:

- The specific technical question or comparison it asks for.
- The languages, techniques, or mechanisms explicitly named.
- Any implicit expectation (e.g. code examples, explanation of
→ underlying
mechanism, comparison table).

Preserve the complete meaning of [INPUT PROMPT]. Do not remove
requirements. Do not introduce new technical requirements beyond what
[INPUT PROMPT] asks. Do not expand the scope beyond the specific
→ question
posed.

Task 2 -- Section-Level Framing (STRICT):

The generated master prompt must explicitly instruct the target
→ language
model that its output is ONE SECTION of an existing chapter, not a
complete report. Specifically, the master prompt must require the
→ target

model to:

- Omit any standalone introduction, motivation, or "in this report we
↪ will"
framing -- begin directly with the technical content.
- Omit any standalone conclusion, summary of the whole chapter, or
forward-looking remarks about other sections.
- Omit its own bibliography or reference list -- citations, if any,
↪ are
consolidated once at the end of the full chapter, not per section.
- Assume the reader already has the context of the broader comparison
(C++ templates vs Java generics vs C# generics) and does not need
↪ it
re-introduced.

Task 3 -- Content Constraints:

Ensure the generated master prompt instructs the target language

↪ model to:

- Maintain a formal, objective, and precise technical writing style
consistent with programming language theory literature.
- Use precise terminology (instantiation, monomorphization, type
↪ erasure,
reification, substitution failure, bounded polymorphism) only where
relevant to [INPUT PROMPT].
- Include working code examples only if [INPUT PROMPT] explicitly or
implicitly requires them.
- Distinguish clearly between compile-time and runtime mechanisms
↪ when
relevant to the specific question.
- Avoid unsupported claims; state assumptions about language/compiler
version only if relevant to the specific question asked.
- ENFORCE STRICT SCOPE CONTROL: address only what [INPUT PROMPT]
↪ asks.
Do not digress into memory footprint, hardware performance
↪ benchmarking,
code bloat, garbage collection overhead, or unrelated language
comparisons unless [INPUT PROMPT] explicitly requires it.
- ENFORCE CONCISENESS: the section must be as short as fully
↪ answering
[INPUT PROMPT] allows. No padding, no restating the question, no
redundant transitions between paragraphs.

Writing Style Constraint (STRICT):

Do not execute [INPUT PROMPT] directly. Do not generate the final

↪ section

content. Do not answer the original technical question. The output
↪ must
contain only the generated master prompt. Do not include
↪ explanations,
comments, editorial notes, placeholders, analysis of the
↪ transformation,
or meta-commentary addressed to the user. Do not use first-person
↪ forms
("I", "we") in the generated master prompt unless explicitly required
↪ by
[INPUT PROMPT].

Output Structure:

Return only the generated master prompt, formatted as ready-to-use
↪ text
and organized using the following headings: System Persona, Goal,
Requirements, Writing Style Constraint, Output Structure. No
↪ additional
text must appear before or after the generated master prompt.

Chapter 2

Task 1 – Language Implementation of Parametric Polymorphism

This chapter groups the prompts used to generate the report content addressing Task 1: *“Explain how each language implements parametric polymorphism (C++ templates, Java generics, C# generics), including when instantiation happens, what exists at runtime, and how this affects performance, type safety, and reflection.”* Each section below merges what were previously three separate, per-language prompts (C++, Java, C#) into a single comparative prompt covering all three languages together, so that a single generation directly produces the cross-language comparison.

Methodological Note on Prompt Construction and Verification:

To ensure the accuracy of the generated analysis, the prompt engineering strategy utilized a decoupled approach. Specifically, individual prompts were tailored for each language (C++, Java, and C#) across five core analytical categories:

1. Timing of instantiation
2. Runtime representation
3. Performance implications
4. Type safety
5. Reflection capabilities

Because multiple distinct queries targeted different facets of the same language, the model’s outputs naturally contained edundant technical information. This redundancy was intentional; it served as the critical first step in the verification process, allowing for the cross-referencing of the model’s overlapping sections and responses to ensure internal consistency before synthesizing the final master prompt.

To simplify the analysis and improve readability within this document, these isolated queries were subsequently merged. The resulting unified master prompt dictates the exact structure, constraints, and content parameters used to generate the comprehensive comparison presented in the following sections.

2.1 Prompt 1: Overview of Parametric Polymorphism

Model: ChatGPT-5.5

Goal: Generate an overview of Parametric Polymorphism and ad-hoc polymorphism, giving details of type erasure compared with generics and monomorphization, and comparing differences between source code, bytecode, and binary code.

Prompt:

- Provide an exhaustive theoretical background on parametric polymorphism.
- Define and thoroughly explain Christopher Strachey’s distinction between parametric polymorphism and ad-hoc polymorphism.
- Detail the two primary implementation strategies discussed in the literature:
 1. Homogeneous translation (type erasure), detailing how it results in one compiled representation for all instantiations.
 2. Heterogeneous translation (reification/monomorphization), detailing how it results in separate code generated per instantiated type.
- Analyze type erasure versus reified generics as the core axis of comparison.
- Analyze compile-time versus runtime characteristics as the second axis of comparison, specifically examining when instantiation happens and what type information survives into the compiled binary or bytecode.
- The entire document must be written in professional, compilation-ready LaTeX code.

Resources:

The generation of the technical report was supported by the following academic and technical references. The theoretical background was derived from foundational works on type systems and polymorphism, including Strachey’s original classification of polymorphism, Cardelli and Wegner’s analysis of type abstraction and polymorphism, and Pierce’s formal treatment of parametric polymorphism and type systems [42, 7, 36, 38]. The analysis of generic implementation strategies, including homogeneous translation, type erasure, heterogeneous translation, reification, and monomorphization, was based on research concerning generic programming implementations and runtime systems, including the .NET Common Language Runtime and Java generic mechanisms [23, 5, 3, 12, 13, 32, 33]. The comparison between compile-time and runtime characteristics was supported by compiler literature and modern programming language implementation documentation [8, 2, 1, 46, 45, 47, 25].

2.2 Prompt 2: Instantiation Time – C++, Java, and C# Compared

Goal: Generate a single, unified, technically rigorous section comparing the *instantiation time* of parametric polymorphism across C++ templates, Java generics, and C# generics, covering compile-time template instantiation, compile-time type erasure, and runtime CLR/JIT-based instantiation within one comparative treatment.

Model: ChatGPT-5.5

Context:

This prompt merges what were previously three separate prompts (one per language) on instantiation time into a single comparative prompt, extending the overview generated in Prompt 1 (Section ??).

System Persona

Act as a world-class computer scientist and principal compiler/runtime engineer specializing in Programming Language Theory (PLT), type systems, and the compiler/virtual-machine execution models of C++, the JVM, and the .NET CLR.

Goal

Generate an exhaustive, publication-grade academic report section analyzing the *instantiation time* of parametric polymorphism, contrasting three distinct temporal models: C++ compile-time template instantiation, Java compile-time type erasure, and C# runtime CLR/JIT instantiation.

Requirements

- **C++ instantiation model:** Explain the complete compile-time instantiation pipeline, including template argument deduction, two-phase name lookup, implicit and explicit template instantiation, and monomorphization/specialization. Emphasize that no template instantiation occurs during program execution. Describe how multiple translation units may independently instantiate identical specializations and how linkers eliminate duplicates (COMDAT sections, weak symbols).
- **Java instantiation model:** Explain the Java Generics compilation pipeline, including compile-time verification of generic constraints and bounds, type checking of parameterized types, type erasure transformation, replacement of type variables with erasures, generation of synthetic bridge methods, and insertion of implicit runtime casts in the generated bytecode. Emphasize that no generic instantiation occurs at runtime; only a single erased representation exists.
- **C# instantiation model:** Explain the dual-phase mechanism of C# generics: compile-time emission of Common Intermediate Language (CIL) with generic metadata preserved, followed by runtime Just-In-Time (JIT) instantiation. Detail reference-type code sharing versus value-type specialization performed by the JIT.

- **Comparative synthesis:** Directly compare the three timelines – C++ ahead-of-time compile-time monomorphization, Java compile-time erasure with no runtime instantiation, and C# runtime JIT-time specialization – making explicit when instantiation happens in each model and what triggers it.
- Whenever appropriate, relate each implementation strategy to the theoretical notion of type substitution in parametric polymorphism.
- **STRICT SCOPE CONTROL:** confine the text strictly to instantiation timing and mechanisms. Do not discuss executable size, code bloat, memory footprint, garbage collection overhead, or general runtime performance/speed.
- Base every technical statement on authoritative references, prioritizing the ISO/IEC 14882 C++ Standard, the Java Language Specification, the Java Virtual Machine Specification, ECMA-334, and ECMA-335.
- Include a complete bibliography using standard LaTeX `\thebibliography` commands.

Writing Style Constraint

- Adopt a formal, objective, and precise scientific writing style comparable to ACM or IEEE publications.
- Use accurate compiler/runtime terminology (e.g., *template argument deduction*, *two-phase lookup*, *monomorphization*, *translation unit*, *COMDAT*, *type erasure*, *bridge method*, *JIT specialization*, *reification*).
- Avoid first-person pronouns and conversational language.
- Provide technically detailed explanations rather than high-level summaries.

Output Structure

The generated document must consist entirely of valid, self-contained, compilation-ready LaTeX code, organized as:

1. Introduction: Three Models of Instantiation Time
2. C++ Templates: Compile-Time Instantiation Pipeline
3. Java Generics: Compile-Time Type Erasure
4. C# Generics: Runtime CLR/JIT Instantiation
5. Comparative Timeline: When Instantiation Happens in Each Language
6. Conclusion

Primary References

- **ISO/IEC JTC 1/SC 22/WG 21, *ISO/IEC 14882:2020/2024 Programming Languages — C++* [20].**

- David Vandevoorde, Nicolai M. Josuttis, and Douglas Gregor, *C++ Templates: The Complete Guide*, 2nd Edition, Addison-Wesley Professional, 2017 [47].
- Bjarne Stroustrup, *The C++ Programming Language*, 4th Edition, Addison-Wesley Professional, 2013 [43].
- Oracle, *The Java Language Specification, Java SE 24 Edition* [32].
- Oracle, *The Java Virtual Machine Specification, Java SE 24 Edition* [33].
- Gilad Bracha, Martin Odersky, David Stoutamire, and Philip Wadler, “Making the Future Safe for the Past: Adding Genericity to the Java Programming Language,” OOPSLA 1998 [5].
- Gilad Bracha, *Generics in the Java Programming Language*, Sun Microsystems, 2004 [3].
- Atsushi Igarashi, Benjamin C. Pierce, and Philip Wadler, “Featherweight Java: A Minimal Core Calculus for Java and GJ,” TOPLAS 2001 [17].
- Ecma International, *ECMA-334: C# Language Specification*, 2023 [12].
- Ecma International, *ECMA-335: Common Language Infrastructure (CLI)*, 2024 [13].
- Andrew Kennedy and Don Syme, “Design and Implementation of Generics for the .NET Common Language Runtime,” PLDI, ACM, 2001 [23].
- Jeffrey Richter, *CLR via C#*, 4th Edition, Microsoft Press, 2012 [40].
- Benjamin C. Pierce, *Types and Programming Languages*, MIT Press, 2002 [36].
- Alfred V. Aho, Monica S. Lam, Ravi Sethi, and Jeffrey D. Ullman, *Compilers: Principles, Techniques, and Tools*, 3rd Edition, Pearson, 2022 [1].
- Keith D. Cooper and Linda Torczon, *Engineering a Compiler*, 3rd Edition, Morgan Kaufmann, 2022 [8].

2.3 Prompt 3: Runtime Representation – C++, Java, and C# Compared

Goal: Generate a single, unified section analyzing the *runtime representation* of parametric polymorphism across C++ templates (monomorphized native code, no

generic metadata), Java generics (erased representation), and C# generics (reified CLR representation).

Model: ChatGPT-5.5

System Persona

Act as a world-class computer scientist and principal compiler/runtime engineer specializing in Programming Language Theory (PLT), compiler backends (GCC/Clang), the Java Virtual Machine, and the .NET Common Language Runtime.

Goal

Generate an exhaustive, publication-grade academic report section titled “Runtime Representation of Parametric Polymorphism in C++, Java, and C#”, analyzing what exists in memory and in the executing environment after compilation/loading for each language.

Requirements

- **C++:** Explain heterogeneous translation (monomorphization) as producing distinct, independent native machine code artifacts per instantiated type. Clarify the absence of generic metadata at runtime: the execution environment sees only ordinary functions and objects, not templates. Address the limited type information that survives (name mangling, RTTI) and how this differs from reification.
- **Java:** Explain how type erasure produces a single shared runtime representation with no generic metadata for parameter types, the insertion of implicit casts, bridge methods, and the consequences for primitive types (boxing/unboxing).
- **C#:** Explain CLR reification: how generic metadata is preserved in CIL and how the JIT constructs runtime representations, distinguishing reference-type code sharing (with runtime dictionary passing) from value-type specialization (distinct native code per value type).
- **Comparative synthesis:** Directly contrast the three runtime representations along the axis of metadata presence and code-artifact granularity: C++ (no metadata, per-type native code), Java (shared erased representation, no metadata), C# (preserved metadata, partial code sharing).
- **STRICT SCOPE CONTROL:** confine the text strictly to runtime representation and memory/metadata structure. Do not discuss compilation times, code-bloat metrics, manual memory management, or CPU cache performance.
- Base every technical statement on authoritative references (ISO C++ Standard, JLS, JVMs, ECMA-335, and established compiler/runtime literature).
- Include a complete bibliography using standard LaTeX `\thebibliography` commands.

Writing Style Constraint

- Formal, objective, precise scientific style comparable to ACM/IEEE publications.
- Use accurate terminology: *monomorphization*, *name mangling*, *RTTI*, *type erasure*, *bridge method*, *reification*, *generic dictionary*, *CIL metadata*.
- Avoid first-person pronouns and conversational language.

Output Structure

1. Introduction: Three Models of Runtime Representation
2. C++: Monomorphized Native Code and Absence of Generic Metadata
3. Java: Erased Shared Representation
4. C#: Reified CLR Representation and JIT-Constructed Types
5. Comparative Analysis: Metadata Presence and Code-Artifact Granularity
6. References

Primary References

- David Vandevor, Nicolai M. Josuttis, and Douglas Gregor, *C++ Templates: The Complete Guide*, 2nd Edition, 2017 [47].
- ISO/IEC JTC 1/SC 22/WG 21, *ISO/IEC 14882:2024 Programming Languages — C++* [20].
- Bjarne Stroustrup, *The C++ Programming Language*, 4th Edition, 2013 [43].
- Stanley B. Lippman, *Inside the C++ Object Model*, Addison-Wesley Professional, 1996 [24].
- The Itanium C++ ABI Project, *Itanium C++ ABI Specification* [21].
- Oracle, *The Java Language Specification, Java SE 24 Edition* [32].
- Oracle, *The Java Virtual Machine Specification, Java SE 24 Edition* [33].
- Gilad Bracha, *Generics in the Java Programming Language*, Sun Microsystems, 2004 [3].
- Ecma International, *ECMA-335: Common Language Infrastructure (CLI)*, 2024 [13].
- Andrew Kennedy and Don Syme, “Design and Implementation of Generics for the .NET Common Language Runtime,” PLDI, 2001 [23].
- Jeffrey Richter, *CLR via C#*, 4th Edition, 2012 [40].

2.4 Prompt 4: Performance Characteristics – C++, Java, and C# Compared

Goal: Generate a single, unified section analyzing the runtime performance implications of each language’s implementation strategy, following directly from the runtime representation discussion.

Model: ChatGPT-5.5

System Persona

Act as a Senior Computer Science Researcher and Expert in Programming Language Theory, Compiler Design, JVM Architecture, and .NET CLR Architecture.

Goal

Generate a comprehensive, technically rigorous academic section analyzing and directly comparing the execution performance characteristics of C++ templates (compile-time monomorphization), Java generics (type erasure), and C# generics (runtime reification with JIT specialization).

Requirements

- Establish continuity with the following preceding theoretical context: “Both implementation strategies preserve static type safety, although they achieve this objective differently. In implementations based on type erasure, generic correctness is verified entirely during compilation. Although type parameters disappear from the generated executable, compiler-inserted casts preserve the guarantees established by the static type system. Implementations based on reified generics retain complete type information throughout compilation and generate type-specific implementations directly.”
- **C++:** Detail the performance advantages of monomorphization: static dispatch, elimination of boxing/unboxing, and facilitation of aggressive compiler optimizations such as inlining.
- **Java:** Detail the runtime overheads introduced by type erasure: pointer indirection, mandatory boxing/unboxing for primitive types, impacts on heap allocation and data locality. Discuss JIT-level mitigations (escape analysis, inline caching, devirtualization) only insofar as they mitigate generic-specific overhead.
- **C#:** Detail the distinct performance behavior of value-type specialization (per-type native code, no boxing) versus reference-type sharing (single implementation, dictionary-passed type handles) under JIT compilation.
- **Comparative synthesis:** Directly compare static dispatch/AOT optimization (C++), erasure-induced boxing overhead (Java), and the hybrid value/reference-type performance profile (C#).
- **STRICT SCOPE CONTROL:** confine the analysis to execution performance, runtime throughput, memory layout, and cache behavior. Do not digress

into compilation times, language ergonomics, or general garbage collection mechanics unless directly tied to generic-structure throughput.

- Rely on authoritative sources: ISO C++ Standard, JLS, JVMMS, ECMA-335, C# Language Specification, and established compiler literature. Include consistent citations and a final bibliography.

Writing Style Constraint

- Formal, objective, rigorous scientific style; no first-person pronouns.
- Precise terminology: *static dispatch*, *boxing overhead*, *data locality*, *reification*, *monomorphization*, *type erasure*, *JIT specialization*, *cache lines*.
- Ground all claims in established compiler theory, VM/CLR specifications, and hardware-interaction mechanics; avoid unsupported assertions.

Output Structure

1. Introduction: Performance as a Consequence of Implementation Strategy
2. C++: Static Dispatch and Compile-Time Optimization
3. Java: Type Erasure and Boxing Overhead
4. C#: Value-Type Specialization vs. Reference-Type Sharing
5. Comparative Analysis: Throughput, Memory Layout, and Cache Behavior
6. References

Primary References

- David Vandevoorde, Nicolai M. Josuttis, and Douglas Gregor, *C++ Templates: The Complete Guide*, 2nd Edition, 2017 [47].
- Steven S. Muchnick, *Advanced Compiler Design and Implementation*, Morgan Kaufmann, 1997 [31].
- Oracle, *The Java Language Specification, Java SE 21/24 Edition* [16].
- ECMA International, *ECMA-335: Common Language Infrastructure (CLI)*, 2012/2024 [13].
- Microsoft Corporation, *C# Language Specification*, 2023 [29].

2.5 Prompt 5: Type Safety – C++, Java, and C# Compared

Goal: Generate a single, unified section analyzing static type safety in C++ templates, Java generics, and C# generics, illustrated with explicative code examples for each language.

Model: ChatGPT-5.5

System Persona

Act as an expert academic researcher and computer scientist specializing in programming language theory, compiler design, and cross-language generic type-system semantics.

Goal

Generate an exhaustive, academic-grade LaTeX section analyzing static type safety across C++ templates (reified generics via monomorphization), Java generics (type erasure), and C# generics (reified generics via CLR).

Requirements

- Integrate the following theoretical premise as the foundation: “Both implementation strategies preserve static type safety, although they achieve this objective differently. In implementations based on type erasure, generic correctness is verified entirely during compilation. Although type parameters disappear from the generated executable, compiler-inserted casts preserve the guarantees established by the static type system. Implementations based on reified generics retain complete type information throughout compilation and generate type-specific implementations directly. Since each specialization corresponds to a concrete type, correctness is preserved without requiring erased representations or implicit runtime casts.”
- **C++:** Detail how templates enforce type safety at compile time via monomorphization, retaining complete type information, with a simple explanatory code example.
- **Java:** Detail how type erasure verifies generic correctness entirely at compile time via compiler-inserted casts, with a simple explanatory code example.
- **C#:** Detail how reified generics retain complete type information throughout compilation and generate type-specific implementations, with a simple explanatory code example.
- Relate each code example directly to the theoretical concepts of reification, erasure, and concrete type specialization.
- Rely on authoritative academic or technical sources (ISO/IEC 14882, JLS, ECMA-335, and peer-reviewed PLT literature).
- **STRICT SCOPE CONTROL:** confine the text strictly to type-safety mechanisms and the theoretical strategies of implementation. Do not digress into

memory footprint, hardware performance, code bloat, compilation times, or garbage collection overhead.

Writing Style Constraint

- Formal, objective, scientific writing style; no first-person pronouns.
- Distinguish clearly between theoretical concepts (parametric polymorphism) and implementation mechanisms (monomorphization, erasure, reification).
- Avoid unsupported claims and clearly identify assumptions.

Output Structure

1. Introduction: Type Safety Under Different Implementation Strategies
2. C++: Type Safety via Reified Templates
3. Java: Type Safety via Type Erasure
4. C#: Type Safety via CLR Reification
5. Comparative Synthesis
6. References

Format each code example using the standard `listings` package.

Primary References

- David Vandevorde, Nicolai M. Josuttis, and Douglas Gregor, *C++ Templates: The Complete Guide*, 2nd Edition, 2017 [48].
- Bjarne Stroustrup, *The C++ Programming Language*, 4th Edition, 2013 [44].
- Benjamin C. Pierce, *Types and Programming Languages*, MIT Press, 2002 [37].
- James Gosling et al., *The Java Language Specification* [15].
- Microsoft Corporation, *C# Language Specification* [29].
- ECMA International, *ECMA-335: Common Language Infrastructure (CLI)* [13].

2.6 Prompt 6: Reflection and Runtime Type Information – C++, Java, and C# Compared

Goal: Generate a single, unified section analyzing reflection and runtime type information (RTTI) across C++ templates, Java generics, and C# generics.

Model: ChatGPT-5.5

System Persona

Act as a Senior Research Scientist in Programming Language Theory, Compiler Architecture, and Type Systems, with expertise in C++ RTTI, Java reflection, and `.NET System.Reflection`.

Goal

Generate a rigorous, technical-academic report section analyzing reflection and RTTI capabilities across C++, Java, and C#, grounded in the theoretical premise that runtime availability of generic type information depends directly on the chosen representation strategy: type-erased implementations discard most generic information, while reified implementations preserve it for runtime introspection.

Requirements

- **C++:** Detail the application and behavior of standard RTTI mechanisms (`typeid`, `std::type_info`, `dynamic_cast`) on instantiated template classes/functions, and how monomorphization preserves concrete type information per specialization.
- **Java:** Detail how type erasure constrains reflection: why `instanceof` against parameterized types is disallowed, why `getClass()` yields an identical reference across structurally distinct generic instantiations, and how the type-token pattern (`Class<T>`) is used as a workaround. Include a compilation-ready Java code example.
- **C#:** Detail how CLR reification preserves generic metadata for full runtime introspection via `System.Reflection`, including `IsGenericType` and `GetGenericArguments()`. Include a compilation-ready C# code example.
- **Comparative synthesis:** Directly contrast the three reflection models along the axis of metadata availability: C++ (per-specialization RTTI, no generic-level metadata), Java (single shared `Class` object, no parameter recovery without workarounds), C# (full runtime introspection of type arguments).
- **STRICT SCOPE CONTROL:** confine the text strictly to reflection, metadata availability, and type introspection capabilities. Do not discuss memory footprint, hardware performance, code bloat, or garbage collection overhead.
- Base every technical statement on authoritative references and avoid fabricated citations.

Writing Style Constraint

- Formal, objective, highly scientific academic tone; no first-person pronouns.
- Precise terminology: *RTTI*, *type erasure*, *reification*, *type token*, *monomorphization*, *generic metadata*.
- Logical progression from theoretical foundations to technical analysis.

Output Structure

1. Introduction: Reflection as a Function of Representation Strategy
2. C++: RTTI on Monomorphized Templates
3. Java: Type Erasure and Reflection Limitations
4. C#: Full Runtime Introspection via CLR Reification
5. Comparative Analysis
6. References

Primary References

- Bjarne Stroustrup, *The C++ Programming Language*, 4th Edition, 2013 [44].
- David Vandevoorde, Nicolai M. Josuttis, and Douglas Gregor, *C++ Templates: The Complete Guide*, 2nd Edition, 2017 [49].
- cppreference.com, *C++ Language Reference: RTTI and typeid* [9].
- Oracle Corporation, *The Java Language Specification* [22].
- Oracle Corporation, *The Java Virtual Machine Specification* [35].
- Jeffrey Richter, *CLR via C#: A Developer's Tour of the Common Language Runtime*, 4th Edition, 2012 [41].

Chapter 3

Task 2 – Metaprogramming Across C++, Java, and C#

This chapter groups the prompts used to generate the report content addressing Task 2: *“Provide examples of metaprogramming in C++ using templates (e.g. compile-time computation, SFINAE / type traits) and compare them with what is and isn’t possible using Java and C# generics, illustrating the practical limits of metaprogramming in each ecosystem.”* The sections retain their existing individual structure – C++-native techniques are treated first, followed by dedicated sections analyzing the limits of Java and of C#, and finally the workarounds used in Java and C# to approximate C++ metaprogramming.

3.1 Prompt 1: Concrete C++ Template Examples of Compile-Time Computation

Goal: Generate a detailed academic section of compile-time computation in C++ templates. For each example, explain what the compiler actually generates and why the computation happens entirely before runtime. Include full compilable code.

Model: ChatGPT-5.5

Prompt:

System Persona:

Act as an expert in programming language theory and advanced C++
↪ software engineering.

Goal:

Write a single, highly focused section of a larger technical-academic
↪ chapter discussing parametric polymorphism and metaprogramming
↪ across C++, Java, and C#. This specific section must strictly
↪ focus on demonstrating and explaining C++ compile-time
↪ computation mechanisms.

Requirements:

Task 1 -- Technical Content:

- * Provide 2-3 concrete, fully compilable C++ examples of compile-time
→ computation.
- * At least one example must demonstrate recursive compile-time
→ computation using traditional C++ templates (e.g., calculating a
→ factorial or Fibonacci sequence).
- * At least one example must demonstrate compile-time computation
→ using `constexpr` functions.
- * For each example, provide a precise technical explanation of what
→ the compiler actually generates during compilation.
- * For each example, explicitly detail the language mechanisms and
→ compiler rules that guarantee the computation executes entirely
→ before runtime.

Task 2 -- Section-Level Framing (STRICT):

- * Treat your output strictly as ONE SECTION of an existing, larger
→ chapter.
- * Assume the reader already possesses the context of the broader
→ comparison between C++, Java, and C\# generics/templates; do not
→ re-introduce this context.
- * Omit any standalone introduction, motivation, or conversational
→ framing (e.g., do not use phrases like "In this section, we will
→ examine..."). Begin immediately with the technical exposition.
- * Omit any standalone conclusion, summary of the chapter, or
→ forward-looking remarks.
- * Omit citations, bibliographies, or reference lists.

Task 3 -- Content Constraints:

- * ENFORCE STRICT SCOPE CONTROL: Address only the requirements stated
→ above. Do not digress into discussions of memory footprint,
→ hardware performance benchmarking, code bloat, garbage
→ collection, or unrelated language comparisons.
- * Distinguish clearly between compile-time evaluation (e.g., template
→ instantiation, constant expression evaluation) and runtime
→ execution.
- * Use precise terminology (e.g., template instantiation, constant
→ evaluation, monomorphization, substitution) strictly where
→ relevant.
- * State compiler or language standard assumptions (e.g., C++11,
→ C++14, C++20) only if strictly required to contextualize the
→ `constexpr` or template mechanisms used in your examples.

* ENFORCE CONCISENESS: The section must be as concise as possible
→ while fully satisfying the technical requirements. Eliminate all
→ padding, restatements of the prompt, and redundant paragraph
→ transitions.

Writing Style Constraint:

Maintain a formal, objective, and precise technical writing style
→ consistent with advanced programming language theory literature.
→ Do not use first-person forms ("I", "we", "our"). Do not include
→ explanations, comments, editorial notes, placeholders, or
→ meta-commentary addressed to the user.

Output Structure:

Output only the raw Markdown content for this specific section. Use
→ appropriate subsection headings (e.g., `###` or `####`) to
→ organize the examples and explanations. Do not include a
→ top-level title for the chapter. Return exactly the text of the
→ section and nothing else.

Sources:

- The description of **template instantiation**, recursive template specialization, and compile-time computation through template-based metaprogramming is based on [50] and the ISO C++ language specification [18].
- The explanation of **constexpr** functions, constant expression evaluation, and the requirement that **constexpr** expressions must be evaluated during translation is derived from the C++ standard [18] and the technical documentation provided by [10].
- The discussion of general C++ template mechanisms, including template arguments, specialization, and the template instantiation process, is supported by [11] and [50].

3.2 Prompt 2: SFINAE and Type Traits

Goal: Generate a detailed academic section of a C++ example using SFINAE that selects a different function overload depending on whether a type is integral vs. floating point vs. a custom class. Also show the equivalent using C++20 concepts. Explain step by step what "substitution failure" means here and why it isn't a compile error.

Model: ChatGPT-5.5

Prompt:

System Persona:

Act as an expert technical writer and programming language theorist
→ specialized in C++ metaprogramming and compile-time polymorphism.

Goal:

Generate a single, highly focused section for a larger academic and
→ technical chapter on parametric polymorphism and metaprogramming.
→ This section must exclusively demonstrate and explain C++ SFINAE
→ and C++20 Concepts for overload resolution based on type
→ properties.

Requirements:

* Technical Content:

- * Provide a working C++ code example using SFINAE (specifically via
→ `std::enable_if``) that successfully selects a different function
→ overload depending on whether a type is an integral type, a
→ floating-point type, or a custom class.
- * Provide the equivalent working C++ code example utilizing C++20
→ Concepts.
- * Provide a step-by-step explanation of what "substitution failure"
→ means in this specific context and explicitly detail the
→ mechanism of why it does not result in a hard compile error.

* Section-Level Framing (STRICT):

- * Begin directly with the technical content and code. Omit any
→ standalone introduction, motivation, or "in this section we will"
→ framing.
- * Omit any standalone conclusion, summary of the chapter, or
→ forward-looking remarks.
- * Omit any bibliography or reference list; assume citations are
→ handled globally at the end of the full chapter.
- * Assume the reader already has the broader context of parametric
→ polymorphism and does not need general template concepts
→ re-introduced.

* Content Constraints:

- * Distinguish clearly that these are compile-time mechanisms.
- * Enforce strict scope control: address only the SFINAE
→ implementation, the C++20 Concepts equivalent, and the
→ explanation of substitution failure. Do not digress into
→ compilation times, memory footprint, code bloat, hardware
→ performance, or other unrelated language comparisons.

- * Enforce conciseness: the section must be as short as fully
 - ↪ answering the technical requirements allows. No padding, no
 - ↪ restating the question, and no redundant transitions between
 - ↪ paragraphs.

Writing Style Constraint:

- Maintain a formal, objective, and precise technical writing style
- ↪ consistent with programming language theory literature. Use
 - ↪ precise terminology (e.g., substitution failure, overload
 - ↪ resolution set, template instantiation) appropriately. Do not
 - ↪ include conversational filler, meta-commentary, explanations of
 - ↪ your process, or first-person forms ("I", "we").

Output Structure:

- Return only the generated section content, formatted in Markdown. Use
- ↪ appropriate heading levels (e.g., `###` or `####`) suitable for
 - ↪ an internal subsection. Use standard Markdown code blocks for the
 - ↪ C++ code examples.

Sources:

- **Stroustrup [44]**: used for the description of template-based compile-time polymorphism, overload resolution involving function templates, and the general role of template substitution mechanisms in C++.
- **cppreference.com [9]**: used for the precise behavior of `std::enable_if`, type traits such as `std::is_integral` and `std::is_floating_point`, and the C++20 Concepts syntax based on constrained templates.
- **ISO/IEC 14882:2020 [19]**: used as the normative reference for SFINAE rules, template argument substitution, immediate context rules, overload resolution, and the C++20 constraint system.

3.3 Prompt 3: Variadic Templates, Template Template Parameters, and CRTP

Goal: Generate a detailed academic section of other core C++ template metaprogramming techniques: variadic templates, template template parameters, and CRTP (Curiously Recurring Template Pattern). For each, give a short example and explain what class of problems it solves that plain function/class templates cannot.

Prompt:

System Persona

Act as a technical-academic writer specialized in programming
→ language theory and advanced C++ software engineering.

Goal

Write a single, highly focused section for a larger chapter on
→ parametric polymorphism and metaprogramming. The section must
→ explain three specific advanced C++ template metaprogramming
→ techniques: variadic templates, template template parameters, and
→ the Curiously Recurring Template Pattern (CRTP).

Requirements

- * **Contextual Integration:** Treat your output as an internal
→ section of an existing chapter. Assume the reader already
→ understands the broader context of parametric polymorphism across
→ C++, Java, and C\#. Do not re-introduce basic generic programming
→ concepts.
- * **Strict Framing:** Begin directly with the technical explanation.
→ Omit any standalone introduction, motivation, or transitional
→ phrasing (e.g., "In this section, we will explore..."). Omit any
→ conclusion, summary, forward-looking remarks, or bibliography.
- * **Technical Scope:** Focus exclusively on variadic templates,
→ template template parameters, and CRTP. Explicitly exclude
→ discussions of SFINAE, type traits, and constexpr, as these are
→ assumed to be covered elsewhere.
- * **Content Execution:** For each of the three covered techniques:
 - * Provide a short, syntactically correct C++ code example.
 - * Explain the specific class of problems the technique solves that
→ plain function or class templates cannot resolve.
- * Clearly distinguish compile-time mechanisms and state relevant
→ assumptions regarding C++ language/compiler versions (e.g., C++11
→ for variadic templates) where necessary.

- * **Strict Scope Control:** Do not digress into memory footprint,
→ code bloat, hardware performance benchmarking, or unrelated
→ language comparisons. Address only the mechanisms requested.

Writing Style Constraint

Maintain a formal, objective, and precise technical writing style

- ↪ consistent with programming language theory literature. Use exact
- ↪ terminology (e.g., instantiation, monomorphization, static
- ↪ dispatch) accurately. Enforce extreme conciseness: the section
- ↪ must be as short as possible while fully answering the prompt.
- ↪ Eliminate all padding, restatements of the core objective, and
- ↪ redundant transitions between paragraphs. Do not use first-person
- ↪ pronouns ("I", "we").

Output Structure

Return only the Markdown-formatted text for the section content. Do

- ↪ not include any meta-commentary, title placeholders, or
- ↪ conversational filler before or after the section text.

Sources:

- [48] provides the main theoretical reference for advanced C++ template metaprogramming, including variadic templates and parameter packs introduced in C++11, template instantiation mechanisms, template template parameters, and advanced template techniques such as the Curiously Recurring Template Pattern (CRTP).
- [9] provides a formal and up-to-date reference for C++ language constructs, including template syntax, compiler behavior during template specialization, and standard-defined semantics of template-related mechanisms.
- [26] is used as a complementary reference for the practical description of C++ templates, including syntax examples and explanations of variadic templates and template parameter mechanisms.

3.4 Prompt 4: Limits of Java Generics in Replicating C++ SFINAE and Type Traits

Goal: Generate a detailed academic section on the limits of Java generics in replicating C++ SFINAE and type traits. Show what is achievable (e.g. bounded type parameters like `<T extends Number>`) and explain precisely why full SFINAE-style compile-time branching on the generic type `T` is impossible in Java, tying the explanation to type erasure.

Model: ChatGPT-5.5

Prompt:

System Persona:

Act as an expert in programming language theory and software

- ↪ engineering, specializing in the comparative analysis of
- ↪ parametric polymorphism and metaprogramming across C++, Java, and
- ↪ C\#.

Goal:

Write a single, concise section of an academic report chapter

- evaluating the limits of Java generics when attempting to
- replicate C++ SFINAE and type-traits behavior. Specifically,
- demonstrate the extent of what is achievable in Java (e.g.,
- bounded type parameters like `<T Number extends>`) and explain
- precisely why full SFINAE-style compile-time branching on a
- generic type `T` is impossible in Java, explicitly tying this
- architectural limitation to Java's implementation of type
- erasure.

Requirements:

- * Section-Level Framing (STRICT): This text is ONE SECTION of an
 - existing, larger chapter. Do NOT write a standalone introduction,
 - motivation, or "in this section we will" framing. Do NOT include
 - a concluding summary, forward-looking remarks, or a bibliography.
 - Assume the reader is already familiar with the broader comparison
 - between C++ templates and Java/C# generics. Begin directly with
 - the technical analysis.
- * Technical Demonstration: Provide a concise, valid Java code example
 - demonstrating the maximum achievable compile-time constraint
 - mechanism in Java, specifically bounded polymorphism (e.g., `<T`
 - `Number extends>`).
- * Theoretical Explanation: Detail exactly why full compile-time
 - branching on generic types (the equivalent of C++ Substitution
 - Failure Is Not An Error) fails in Java. Connect this limitation
 - mechanically and theoretically to type erasure.
- * Terminology: Use precise technical terminology appropriate for
 - programming language theory (e.g., type erasure, bounded
 - polymorphism, monomorphization, compile-time resolution, bounds
 - checking) where directly relevant.
- * Scope Control: Address only the replication of SFINAE/type-traits
 - in Java. Do not digress into memory footprint, hardware
 - performance, C# mechanisms, code bloat, garbage collection, or
 - unrelated language comparisons.

Writing Style Constraint:

- * Maintain a formal, objective, and highly precise academic writing
 - style.
- * Enforce strict conciseness: the section must be as short as fully
 - answering the prompt allows. Avoid padding, restating the
 - question, or using redundant transitions.

- * Clearly distinguish between compile-time mechanisms (bounds checking) and runtime realities (erasure) without confusing the two.
- * Do not use first-person pronouns (e.g., "I", "we", "our").

Output Structure:

Output only the raw technical text of the section. Use standard
 ↳ Markdown for formatting and code blocks. Do not include any
 ↳ overarching titles, placeholders, meta-commentary, or structural
 ↳ elements outside of what is strictly necessary to present the
 ↳ required technical content.

Sources:

- [6] provides the theoretical foundation of Java Generics, including the type erasure implementation model, the separation between compile-time generic type checking and runtime representation, and the limitations of Java parametric polymorphism with respect to type-dependent specialization.
- [14] provides the formal specification of Java bounded type parameters, including the semantics of `extends` bounds, compile-time type checking rules for generic classes and methods, and the erasure process applied to generic type parameters.
- [48] provides the theoretical background for C++ template metaprogramming, including template substitution, SFINAE-based overload resolution, type traits, and compile-time specialization mechanisms used as a comparison with the limitations of Java Generics.

3.5 Prompt 5: Limits of C# Generics in Replicating C++ SFINAE and Type Traits

Goal: Generate a detailed academic section on the limits of C# generics in replicating C++ SFINAE and type traits. Show what's achievable using generic constraints (`where T : ...`) and, if relevant, .NET generic math / static abstract interface members. Explain why C# gets closer to C++ than Java does, but still can't do arbitrary compile-time computation or true SFINAE, tying it to reification vs. monomorphization.

Model: ChatGPT-5.5

Prompt:

System Persona:

Act as an expert academic writer and programming language theorist
 ↳ specializing in the comparative analysis of parametric
 ↳ polymorphism and metaprogramming paradigms.

Goal:

Generate a single, concise, and highly technical section for a larger
→ academic chapter. This section must specifically analyze how C\#
→ approximates C++ SFINAE and type-traits using generic constraints
→ and generic math, while explaining its theoretical limitations.

Requirements:

- * **Section-Level Framing (STRICT):** This text is an intermediate
→ section within an existing chapter. Do NOT include an
→ introduction, motivation, or structural framing (e.g., "In this
→ section, we will..."). Do NOT include a conclusion, summary, or
→ forward-looking transition. Do NOT include a bibliography or list
→ of references. Assume the reader already possesses full context
→ regarding the fundamental differences between C++ templates, Java
→ generics, and C\# generics.
- * **Technical Execution:**
- * Replicate the conceptual equivalent of a C++ SFINAE/type-traits
→ example using C\#.
- * Demonstrate achievable compile-time/type-check-time constraints
→ using C\# bounded polymorphism (``where T : ...``).
- * Illustrate the application of .NET generic math and static abstract
→ interface members to achieve trait-like capabilities.

- * **Theoretical Analysis:**
- * Explain why these C\# mechanisms allow it to approach C++
→ capabilities much more closely than Java does.
- * Explain why C\# inherently cannot support true SFINAE (Substitution
→ Failure Is Not An Error) or arbitrary compile-time computation.
- * Directly tie this limitation to the architectural differences
→ between .NET's runtime reification of generics and C++'s
→ compile-time monomorphization and AST-level substitution.

Writing Style Constraint:

- * Maintain a formal, objective, and precise academic tone consistent
→ with programming language theory literature.
- * Use precise terminology (e.g., bounded polymorphism, reification,
→ monomorphization, substitution failure) strictly where
→ applicable.
- * Clearly distinguish between compile-time validation mechanisms and
→ runtime behavior.

- * Ensure strict scope control: address solely the requested C\#
 - ↪ replication of SFINAE/traits, generic constraints, generic math,
 - ↪ and the reification vs. monomorphization dichotomy. Do not
 - ↪ digress into memory footprint, code bloat, execution performance,
 - ↪ garbage collection, or unrelated language features.
- * Enforce absolute conciseness: provide only the information required
 - ↪ to fulfill the prompt. Omit conversational padding, restatements
 - ↪ of the prompt, and redundant transitional clauses.

Output Structure:

Provide only the finalized text and code blocks for this specific section. Do not output overarching titles, introductory meta-commentary, conversational filler, or any text outside the bounds of the academic section itself.

Sources:

- [28] The Microsoft documentation on C# generics is used for the description of generic constraints, bounded polymorphism through the ‘where’ clause, and the compile-time validation performed by the C# type system when generic arguments are instantiated.
- [27] The Microsoft documentation on .NET Generic Math is used for the analysis of static abstract interface members and interfaces such as ‘INumber<T>’, which provide trait-like abstractions for expressing type capabilities in generic algorithms.
- [48] The book *C++ Templates: The Complete Guide* is used for the comparison with C++ template metaprogramming, specifically regarding SFINAE, type traits, substitution failure, and compile-time template instantiation through monomorphization.
- [13] The ECMA-335 specification is used for the description of the CLR execution model, including the representation of generic types through runtime reification and the role of metadata in preserving generic type information.

3.6 Prompt 6: Instantiation-Time vs. Declaration-Time Checking

Goal: Generate a detailed academic section explaining how the difference between C++ templates (checked only at instantiation) and Java/C# generics (checked at declaration) directly causes specific metaprogramming techniques to be impossible in Java/C#.

Model: ChatGPT-5.5

Prompt:

System Persona:

Act as a technical author and programming language theorist writing a
→ highly specialized section for an academic chapter on programming
→ language theory and software engineering.

Goal:

Generate a single, self-contained section of a larger chapter that
→ explains how the difference between C++ templates (checked only
→ at instantiation) and Java/C\# generics (checked at declaration)
→ directly causes specific metaprogramming techniques to be
→ impossible in Java/C\#. You must detail at least one concrete
→ technique that fails specifically due to this difference.

Requirements:

- * Focus strictly on the mechanisms of type checking:
 - instantiation-time (C++) versus declaration-time (Java/C\#).
 - * Explain the direct causal relationship between these checking mechanisms and the limitations placed on metaprogramming in Java and C\#.
 -
 - * Detail at least one concrete metaprogramming technique (e.g., SFINAE, or structural/duck typing within templates) that is possible in C++ but impossible in Java/C\# because of this specific constraint.
 -
 - * Include concise, precise code examples to illustrate the technique
 - in C++ and explicitly demonstrate why declaration-time checking prevents its equivalent in Java/C\#.
 - * Section-Level Framing (STRICT): Your output is strictly ONE SECTION
 - of an existing, larger chapter.
 - * Omit any standalone introduction, motivation, or "in this section we will" framing. Begin directly with the technical content.
 - * Omit any concluding paragraphs, chapter summaries, or forward-looking transitions.
 - * Omit a bibliography or reference list. Assume citations are handled globally at the end of the full chapter.
 - * Assume the reader already has the broader context of C++ templates versus Java/C\# generics. Do not re-introduce basic concepts.
-
- * Content Constraints:
 - * Use precise terminology (e.g., instantiation, substitution failure, bounded polymorphism, type erasure, reification) where strictly relevant to the prompt.
 - * Clearly distinguish between compile-time and runtime mechanisms as they relate to these type-checking models.

- * Enforce strict scope control. Do not digress into memory footprint,
 - ↪ hardware performance benchmarking, code bloat, garbage collection
 - ↪ overhead, or unrelated language comparisons. Address only the
 - ↪ type-checking differences and the resulting metaprogramming
 - ↪ limitations.

Writing Style Constraint:

- * Maintain a formal, objective, and precise technical writing style
 - ↪ consistent with programming language theory literature.
- * Enforce absolute conciseness: the section must be exactly as short
 - ↪ as necessary to fully answer the prompt. Include no padding, do
 - ↪ not restate the prompt, and avoid redundant transitions between
 - ↪ paragraphs.
- * Avoid unsupported claims; state assumptions about language or
 - ↪ compiler versions only if strictly necessary to explain the
 - ↪ technique.

Output Structure:

- Return only the raw text of the generated section, formatted using
- ↪ standard Markdown. Do not include any meta-commentary,
 - ↪ introductory remarks, conversational filler, or placeholders.
 - ↪ Begin immediately with the first sentence of the section's
 - ↪ technical content.

Sources:

- [48] This reference is used as the primary source for the C++ template instantiation model, including delayed type checking of dependent expressions, template substitution, and SFINAE-based overload resolution. The description of compile-time capability detection through substitution failure and structural requirements in templates is based on the template metaprogramming mechanisms discussed in this work.
- [4] This reference provides the theoretical background for Java generics, particularly the declaration-time type checking model, bounded polymorphism, and the requirement that all operations performed on a generic type parameter must be justified by its declared bounds. It supports the comparison between Java generic constraints and the structural flexibility provided by C++ template substitution.
- [13] This reference is used for the C# generic type system through the Common Language Infrastructure specification. It provides the formal description of generic parameters, constraints, and compile-time verification rules that restrict generic code to operations guaranteed by declared constraints rather than properties discovered during instantiation.

3.7 Prompt 7: External Metaprogramming Workarounds in Java and C#

Goal: Generate a detailed academic section on the external metaprogramming workarounds in Java and C#. Explain why these are workarounds external to the generics system rather than native metaprogramming.

Model: ChatGPT-5.5

Prompt:

System Persona:

Act as an academic technical writer specializing in programming
→ language theory and software architecture.

Goal:

Write exactly ONE SECTION of an existing comprehensive chapter on
→ parametric polymorphism and metaprogramming. This specific
→ section must describe the common workarounds used in Java and C#
→ to simulate the native metaprogramming capabilities of C++
→ templates, and architecturally explain why these approaches are
→ external to their respective generics systems rather than native
→ metaprogramming.

Requirements:

- * Task 1: Technical Content Specification
- * Detail Java's common workarounds: the use of runtime reflection,
→ compile-time annotation processors, and code generation utilities
→ (e.g., Lombok).
- * Detail C#'s common workarounds: the use of source generators and
→ T4 templates.
- * Articulate the structural and theoretical reasons why these
→ mechanisms are considered external workarounds layered over or
→ completely separate from the Java and C# generics systems,
→ contrasting them with the native, intrinsic compile-time
→ evaluation of C++ templates.
- * Clearly differentiate between mechanisms that operate at runtime
→ versus those that operate at compile-time or pre-compilation.

- * Task 2: Section-Level Framing
- * Output ONLY the body text for this specific section.
- * Omit any standalone introduction, motivation, or prefatory remarks.
→ Do not use phrases such as "In this section we will examine...".
- * Omit any conclusion, summary, or forward-looking remarks to
→ subsequent sections.

- * Omit bibliographies or reference lists entirely.
- * Assume the reader already possesses full theoretical context
 - ↪ regarding C++ templates, Java type erasure, and C\# reified
 - ↪ generics. Begin directly with the technical exposition of the
 - ↪ workarounds.

- * Task 3: Content Constraints
- * Maintain a formal, objective, and precise technical writing style
 - ↪ consistent with programming language theory literature.
- * Use accurate terminology (e.g., compile-time, runtime, reflection,
 - ↪ AST manipulation, macro expansion) strictly where relevant.
- * Include brief, illustrative code snippets only if strictly
 - ↪ necessary to demonstrate a specific structural workaround
 - ↪ mechanism; otherwise, rely on rigorous theoretical explanation.
- * Strictly enforce scope control: address only the workarounds and
 - ↪ their relationship to the generics system. Do not digress into
 - ↪ memory footprint, hardware performance benchmarking, garbage
 - ↪ collection overhead, or unrelated architectural comparisons.
- * Enforce absolute conciseness: ensure the section is as short as
 - ↪ comprehensively answering the prompt allows. Avoid padding,
 - ↪ redundant paragraph transitions, or restating the prompt's
 - ↪ premises.

Writing Style Constraint:

Do not include any meta-commentary, explanations, placeholders, or

- ↪ out-of-character text. Do not use first-person pronouns ("I",
- ↪ "we"). Produce only the final text of the requested section.

Output Structure:

Return the technical text of the section directly. Use standard

- ↪ Markdown for logical structuring (such as subheadings for
- ↪ language-specific mechanisms), but do not include a top-level
- ↪ chapter title.

Sources:

- [34] is used as a reference for: Java annotation processing, compiler-based analysis of source elements, and compile-time generation of additional source files. The same reference is also used to describe Java reflection as a runtime mechanism based on class metadata inspection.
- [39] is used as a reference for: Java code generation through compiler extensions, source transformation techniques, and the generation of program elements

outside the Java generics mechanism.

- [30] is used as a reference for: C# Source Generators based on Roslyn, T4 Text Templates, and their role as external code generation mechanisms separated from the C# generic type system.
- [50] is used as a reference for: the C++ template instantiation model, compile-time substitution of template parameters, and the integration of metaprogramming capabilities directly into the language construct.

Bibliography

- [1] Alfred V. Aho, Monica S. Lam, Ravi Sethi, and Jeffrey D. Ullman. *Compilers: Principles, Techniques, and Tools*. Pearson, 3 edition, 2022.
- [2] Andrew W. Appel. *Modern Compiler Implementation in C*. Cambridge University Press, 1998.
- [3] Gilad Bracha. Generics in the java programming language, 2004.
- [4] Gilad Bracha. Generics in the java programming language. Technical report, Sun Microsystems, 2004.
- [5] Gilad Bracha, Martin Odersky, David Stoutamire, and Philip Wadler. Making the future safe for the past: Adding genericity to the java programming language. In *Proceedings of the 13th ACM SIGPLAN Conference on Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA)*, pages 183–200. ACM, 1998.
- [6] Gilad Bracha, Martin Odersky, David Stoutamire, and Philip Wadler. Making the future safe for the past: Adding genericity to the java programming language. In *Proceedings of the ACM Conference on Object-Oriented Programming Systems, Languages and Applications (OOPSLA)*. ACM, 1998.
- [7] Luca Cardelli and Peter Wegner. On understanding types, data abstraction, and polymorphism. *ACM Computing Surveys*, 17(4):471–523, 1985.
- [8] Keith D. Cooper and Linda Torczon. *Engineering a Compiler*. Morgan Kaufmann, 3 edition, 2022.
- [9] cppreference.com. *C++ Language Reference: RTTI and typeid*. cppreference.com, 2025. Online reference for C++ RTTI mechanisms, `std::type_info`, `typeid`, and `dynamic_cast`.
- [10] cppreference.com. `constexpr` specifier. <https://en.cppreference.com/w/cpp/language/constexpr>, 2025.
- [11] cppreference.com. Templates. <https://en.cppreference.com/w/cpp/language/templates>, 2025.
- [12] Ecma International. Ecma-334: C# language specification, 2023.
- [13] Ecma International. Ecma-335: Common language infrastructure (cli), 2024.

- [14] James Gosling, Bill Joy, Guy Steele, Gilad Bracha, and Alex Buckley. *The Java Language Specification*. Addison-Wesley, 8th edition, 2014.
- [15] James Gosling, Bill Joy, Guy Steele, Gilad Bracha, and Alex Buckley. *The Java Language Specification, Java SE 8 Edition*. Oracle, 2014.
- [16] James Gosling, Bill Joy, Guy Steele, Gilad Bracha, and Alex Buckley. *The Java Language Specification*. Oracle, java se 21 edition, 2023.
- [17] Atsushi Igarashi, Benjamin C. Pierce, and Philip Wadler. Featherweight java: A minimal core calculus for java and gj. *ACM Transactions on Programming Languages and Systems*, 23(3):396–450, 2001.
- [18] ISO/IEC. *ISO/IEC 14882:2020 Programming Languages – C++*. International Organization for Standardization, 2020.
- [19] ISO/IEC JTC 1/SC 22/WG 21. Programming languages — c++. Technical Report ISO/IEC 14882:2020, International Organization for Standardization, 2020.
- [20] ISO/IEC JTC 1/SC 22/WG 21. ISO/IEC 14882:2024 Programming Languages — C++, 2024. International Standard.
- [21] Itanium C++ ABI Project. Itanium c++ abi: C++ abi specification, 2024. C++ Application Binary Interface specification including name mangling and RTTI representation.
- [22] James Gosling and Bill Joy and Guy Steele and Gilad Bracha and Alex Buckley. The java language specification. Technical report, Oracle, 2023.
- [23] Andrew Kennedy and Don Syme. Design and implementation of generics for the .net common language runtime. In *Proceedings of the ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*, pages 1–12. ACM, 2001.
- [24] Stanley B. Lippman. *Inside the C++ Object Model*. Addison-Wesley Professional, Reading, MA, 1996.
- [25] John McCall. Generics in swift: Implementation and runtime representation, 2019.
- [26] Microsoft. *C++ Templates Documentation*. Microsoft Learn, 2024.
- [27] Microsoft. Generic math in .net. <https://learn.microsoft.com/en-us/dotnet/standard/generics/math>, 2025.
- [28] Microsoft. Generics in c#. <https://learn.microsoft.com/en-us/dotnet/csharp/fundamentals/types/generics>, 2025.
- [29] Microsoft Corporation. *C# Language Specification*. Microsoft, 2023.

- [30] Microsoft Corporation. Code generation in c#: Source generators and t4 text templates. <https://learn.microsoft.com/dotnet/csharp/roslyn-sdk/source-generators-overview>, 2024.
- [31] Steven S. Muchnick. *Advanced Compiler Design and Implementation*. Morgan Kaufmann, 1997.
- [32] Oracle. The java language specification, java se 24 edition, 2024.
- [33] Oracle. The java virtual machine specification, java se 24 edition, 2024.
- [34] Oracle Corporation. Java compiler api and annotation processing. <https://docs.oracle.com/javase/>, 2024.
- [35] Oracle Corporation. The java virtual machine specification. [<https://docs.oracle.com/javase/specs/jvms/>] (<https://docs.oracle.com/javase/specs/jvms/>), 2025.
- [36] Benjamin C. Pierce. *Types and Programming Languages*. MIT Press, Cambridge, MA, 2002.
- [37] Benjamin C. Pierce. *Types and Programming Languages*. MIT Press, 2002.
- [38] Benjamin C. Pierce, editor. *Advanced Topics in Types and Programming Languages*. MIT Press, Cambridge, MA, 2005.
- [39] Project Lombok. Project lombok: Spice up your java. <https://projectlombok.org/>, 2024.
- [40] Jeffrey Richter. *CLR via C#*. Microsoft Press, 4 edition, 2012.
- [41] Jeffrey Richter. *CLR via C#: A Developer's Tour of the Common Language Runtime*. Microsoft Press, 4th edition, 2012.
- [42] Christopher Strachey. Fundamental concepts in programming languages, 1967. Lecture Notes, International Summer School in Computer Programming, Copenhagen.
- [43] Bjarne Stroustrup. *The C++ Programming Language*. Addison-Wesley Professional, 4 edition, 2013.
- [44] Bjarne Stroustrup. *The C++ Programming Language*. Addison-Wesley Professional, 4th edition, 2013.
- [45] The Rust Project Developers. Monomorphization and code generation, 2024.
- [46] The Rust Project Developers. The rust reference, 2024. Official Rust language specification.
- [47] David Vandevoorde, Nicolai M. Josuttis, and Douglas Gregor. *C++ Templates: The Complete Guide*. Addison-Wesley Professional, 2 edition, 2017.

- [48] David Vandevorde, Nicolai M. Josuttis, and Douglas Gregor. *C++ Templates: The Complete Guide*. Addison-Wesley Professional, 2nd edition, 2017.
- [49] David Vandevorde, Nicolai M. Josuttis, and Douglas Gregor. *C++ Templates: The Complete Guide*. Addison-Wesley Professional, 2nd edition, 2017.
- [50] David Vandevorde, Nicolai M. Josuttis, and Douglas Gregor. *C++ Templates: The Complete Guide*. Addison-Wesley Professional, 2 edition, 2017.